
Babel-formal: Translation of Proofs between Lean and Rocq

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 In this work, we investigate using proof terms (the low-level representation of
2 formal proofs) as a pivot language for translating proof scripts between proof
3 assistants and across tactic sets. Unlike direct proof translation, this approach does
4 not require an aligned training corpus; it only needs aligned context at inference
5 time so that both systems elaborate comparable terms. We compare two strategies:
6 (1) direct script-to-script translation with an off-the-shelf LLM (GPT-5), and (2)
7 proof term translation, where an LLM turns a proof term into a proof script in the
8 target language. We build a small benchmark of aligned sources (14 files, 117
9 lemmas) across Lean and Rocq, and train models to map Lean and Rocq proof
10 terms back to their native proof scripts. Our experiments show that proof term
11 translation works for cross-assistant translation and for translation between tactic
12 sets. It is complementary to using a SoTA off-the-shelf LLM for direct proof script
13 translation (combining both performs best), scales easily in terms of training data,
14 and handles tactic-set translation better (*e.g.*, vanilla Rocq $\rightarrow$ SSReflect).

15 1 Introduction

16 Proof assistants such as Rocq and Lean accumulate over time large libraries of formal mathematics
17 (for example, MathComp [8] for Rocq and Mathlib [2] for Lean), yet proofs are not reusable across
18 systems, even in a strictly aligned context. We study two translation settings. The first is *cross-*
19 *assistant* translation between Lean and Rocq. The second is *cross-tactic* translation inside Rocq,
20 from the standard Rocq tactic language — called here *vanilla Rocq* — to *SSReflect* [6], which is a
21 compact, structured tactic language widely used with MathComp (see an example in Section A.5).

22 Both Rocq and Lean provide a way to pretty print *proof terms* given by the proof assistant (see Fig. 1)
23 in a way that omits details while keeping the structure of a proof. We thus explore using *pretty-printed*
24 *proof terms*, as a pivot language. This removes the need for an aligned training corpus of tactic scripts.
25 At inference time, we only require strictly aligned statements — *i.e.*, same theorem and premises
26 — so that both systems elaborate comparable terms. The approach is related to decompilation in
27 programming languages, where models map assembly to higher-level C programs [13].

28 We compare two strategies:

29 **1. Direct script translation.** An LLM is prompted to convert a proof script directly from source
30 tactics to target tactics. We apply this in both settings (Lean $\leftrightarrow$ Rocq and vanilla Rocq $\rightarrow$ SSReflect),
31 using GPT-5 [10] as the baseline.

32 **2. Proof term translation.** We first extract the proof term, then train models to predict scripts
33 from proof terms (Lean term to Lean script, Rocq term to Rocq script). At inference, we extract the
34 term for the source proof and use the model to produce a script in the target assistant or tactic set




Figure 1: How interactive theorem provers operate

Given a formal statement of a theorem, proving it amounts to constructing a proof term—a precise object that the prover can type-check for correctness. However, these proof terms are typically extremely verbose, capturing every low-level detail, and are impractical to write directly by hand.

To address this, users interact with interactive theorem provers through proof scripts—which are essentially meta-programs designed to generate these proof terms. Over time, this meta-programming layer has evolved in various ways, leading to a diverse ecosystem of proof script languages; often referred to as tactic languages.

(see Fig. 2). For an example of Lean $\rightarrow$ Rocq inference see Section A.6, and Section A.5 for vanilla Rocq $\rightarrow$ SSReflect.

To support this comparison, we build a benchmark of 14 files with 117 lemmas that are strictly aligned between Lean and Rocq. We also mention, but do not pursue, direct low-level term translation between Lean and Rocq [12]. The main drawback of that direction is that the output is not tied to a maintainable proof script.

In a nutshell, proof term translation can be leveraged for translation between proof assistants (*e.g.*, Lean and Rocq) or between different sets of tactics (for example, vanilla Rocq and SSReflect). It is also easy to scale in terms of data, since no aligned training data are needed. Direct tactic translation and proof term translation are complementary, and combining them works best. Code and technical details: anonymous.4open.science/r/babel-formal-25D6/.

2 Methodology

We describe how we build the training data, the inference pipeline, and benchmarks. To build the training data we follow four steps [9]: (1) select a diverse subset of proofs, (2) extract proof terms in Rocq and Lean, (3) generate synthetic backward reasoning traces inferring proof script from proof term, and (4) train models that predict a tactic script from a proof term. Full prompt templates and training details are in Section A.2 and Section A.1.

Proof Subset Selection. Following S1 [9], we create in each case a training set of 1000 entries. For translation between Lean and Rocq, we select a subset of 500 diverse proofs in C-CoRN [3] and 500 diverse proofs in Mathlib. For translation from vanilla Rocq to SSReflect, we select a subset of 1000 entries in MathComp. In both cases, we filter by length, then choose a diverse subset (statements and proofs) using BM25[11].

Proof Term Extraction. For Rocq, we use `coqpyt` [1] to extract, for each proof, its proof term together with the premises and notation declarations. For Lean 4, we use `leanclient` [5] to extract the proof term and the context.

Backward Reasoning Trace Generation. Given each proof term and the corresponding target script, we generate backward reasoning traces with Gemini 2.5 Pro [4]. Backward reasoning trace generation refers here to prompting the LLM with the proof term and the target proof script, and asking it to generate a reasoning process that works backwards from the script. The resulting training dataset follows the structure described in Fig. 3a.

Benchmarks. We evaluate on two benchmarks:

Cross-assistant (Lean $\leftrightarrow$ Rocq). We created 14 files with the same statements in Rocq and Lean, for a total of 117 aligned lemmas. Each lemma is proved natively in both systems, so the proof scripts differ but are type-correct in aligned contexts. From these, we extracted comparable proof terms in both systems. An example of such a source file is in Section A.4.




Figure 2: Training and Inference pipeline

70 *Cross-tactic (vanilla Rocq $\rightarrow$ SSReflect).* We collect 500 Rocq proofs from the C-CoRN library [3],
 71 which is a large mathematical library developed in vanilla Rocq (in contrast to MathComp, which is
 72 built on SSReflect). We use these proofs to assess translation from vanilla Rocq tactics to SSReflect.
 73 To ensure diversity, we pick subsets maximizing document distance (via BM25). An example of
 74 entry is in Section A.5.




Figure 3: Training & Inference strategy

75 **Inference Pipeline.** At inference, we leverage the source proof term and use the appropriate strategy
 76 to obtain a script in the target assistant or tactic set (Fig. 3). The first block in Fig. 3 is initialized
 77 with the goal state and fixed instructions; the full prompt appears in Section A.2.

78 3 Experimental Setup

79 **Models and Training.** We fine-tune Qwen2.5-Coder-32b-Instruct [7] to translate a proof term
 80 into a tactic script. We train two models in two different scenarios: (i) **Babel-translate:** Lean term
 81 $\rightarrow$ Lean script and Rocq term $\rightarrow$ Rocq script. (ii) **Babel-ssreflect:** Rocq term $\rightarrow$ SSReflect script.
 82 These models are used at inference to produce target scripts from source terms (Fig. 2).

83 **Evaluation.** We run verification through Rocq tooling (coq-lsp and Pétanque) [14, 15] and Lean
 84 tooling (leanclient) [5]. We use two generation strategies:

85 *Non-interactive (Babel).* For *Babel-translate* and *Babel-ssreflect*, we sample k candidates per input
 86 and report $\text{pass}@k$ (success if any candidate verifies). No error feedback is provided.

87 *Interactive repair (GPT-5 and Qwen2.5-32b-Coder-Instruct).* For direct tactic translation baselines,
 88 the model receives the proof assistant error on failure and may attempt to repair/continue the proof
 89 for up to 4 rounds. We report $\text{pass}@1$ with feedback, denoted “@1x(4 feedback rounds)” for GPT-5.
 90 For Qwen2.5-32b-Coder-Instruct we additionally sample a small pool per round (e.g., 16 candidates),
 91 denoted “@16x(4 feedback rounds)”, and count success if any candidate in any round verifies
 92 (see Section A.3 for the prompt template)

93 **Tasks.** We evaluate three directions: (i) $\text{Lean} \rightarrow \text{Rocq}$, (ii) $\text{Rocq} \rightarrow \text{Lean}$ (cross-assistant), and
 94 (iii) vanilla $\text{Rocq} \rightarrow \text{SSReflect}$ (cross-tactic), using the benchmarks introduced in Section 2. For
 95 (i) and (ii) we apply *Babel-translate* (proof term $\rightarrow$ native script) and compare with direct GPT-5
 96 translation. For (iii) we apply *Babel-ssreflect* (proof term $\rightarrow$ SSReflect script), again comparing with
 97 direct GPT-5 translation.

98 4 Results and Analysis

Setup	Lean4 $\rightarrow$ Rocq	Rocq $\rightarrow$ Lean4	Rocq $\rightarrow$ SSReflect
GPT-5 (@1x(4 feedbacks))	82.9%	67.5%	16.6%
Babel-translate (@128)	68.3%	40.2%	
Babel-ssreflect (@128)			33.2%
Babel-ssreflect w/o reas. (@128)			21.8%
Babel + GPT-5 combined	89.7%	83.7%	34%
Qwen-2.5-Coder-32b-Instruct (@16x(4 feedbacks))	6.8%	19.7%	0%

Table 1: Benchmark results for different translation setups across three tasks.

99 In the proof script translation baseline, GPT-5 achieves 82.9% success in the direction $\text{Lean} \rightarrow \text{Rocq}$,
 100 and 67.5% in $\text{Rocq} \rightarrow \text{Lean}$. However, on the cross-tactic task vanilla $\text{Rocq} \rightarrow \text{SSReflect}$, it manages
 101 only 16.6% success. In contrast, the cross-tactic translation model attains 33.2% success on $\text{Rocq} \rightarrow$
 102 SSReflect , demonstrating a clear gap in style transfer. In cross-assistant tasks, term translation model
 103 reaches 68.3% on $\text{Lean} \rightarrow \text{Rocq}$ and 40.2% on $\text{Rocq} \rightarrow \text{Lean}$.

104 **Ablation study.** Removing the reasoning component in *Babel-ssreflect* reduces $\text{Rocq} \rightarrow \text{SSReflect}$
 105 from 33.2% to 21.8%. For $\text{Lean} \leftrightarrow \text{Rocq}$, the model trained without reasoning produced mixed
 106 scripts that were not usable, so we do not report scores.

107 **Limitation.** Babel models do not receive interactive feedback during proof generation. They see
 108 the initial goal, the context, and the proof term only, then must produce a complete script. Interactive
 109 repair, as used by GPT-5, may help in some cases.

110 5 Conclusion and Future Work

111 We introduced proof term translation as a way to translate proofs across assistants and tactic styles.
 112 Unlike direct tactic translation, this approach does not need parallel scripts for training. On our
 113 benchmarks it brings clear gains for style transfer and remains competitive for cross-assistant
 114 translation. It is complementary to direct translation, and combining the two gives the best results.

115 **Future work.** We see three directions: (i) **Scale.** Train on larger corpora and more domains, since
 116 the approach does not require aligned scripts. (ii) **Feedback.** Add type checker feedback during
 117 generation. A simple loop that feeds proof assistant errors into the reasoning blocks could guide
 118 local repairs and improve robustness. (iii) **Source-to-source translation ($\text{Lean} \leftrightarrow \text{Rocq}$).** Use
 119 off-the-shelf LLMs for source-to-source translation, and bootstrap proof script translation to check
 120 whether the source translation is correct.

References

- [1] Pedro Carrott, Nuno Saavedra, Kyle Thompson, Sorin Lerner, João F. Ferreira, and Emily First. CoqPy: Proof navigation in python in the era of LLMs. *arXiv preprint arXiv:2405.04282*, 2024.
- [2] Mathlib Community. The lean mathematical library. *arXiv preprint arXiv:1912.09375*, 2019.
- [3] Luís Cruz-Filipe, Herman Geuvers, and Freek Wiedijk. C-corn, the constructive coq repository at nijmegen. In *Mathematical Knowledge Management (MKM 2004)*, volume 3119 of *Lecture Notes in Computer Science*, pages 88–103. Springer, 2004.
- [4] Google DeepMind. Gemini 2.5 Pro: Advanced multimodal model. <https://ai.googleblog.com/2024/07/gemini-25-our-newest-gemini-model-with.html>, 2024.
- [5] Oliver Dressler. Leanclient: Python client to interact with the lean4 language server, 1 2025.
- [6] Georges Gonthier, Assia Mahboubi, and Enrico Tassi. Ssreflect: Small scale reflection in Coq. Coq documentation.
- [7] Binyuan Hui, Jian Yang, Zeyu Cui, Jiayi Yang, Dayiheng Liu, Lei Zhang, Tianyu Liu, Jiajun Zhang, Bowen Yu, Keming Lu, Kai Dang, Yang Fan, Yichang Zhang, An Yang, Rui Men, Fei Huang, Bo Zheng, Yibo Miao, Shanghaoran Quan, Yunlong Feng, Xingzhang Ren, Xuancheng Ren, Jingren Zhou, and Junyang Lin. Qwen2.5-Coder technical report. 2024.
- [8] Assia Mahboubi and Enrico Tassi. Mathematical components. <https://math-comp.github.io>, 2020.
- [9] Niklas Muennighoff, Zitong Yang, Weijia Shi, Xiang Lisa Li, Fei-Fei Li, Hannaneh Hajishirzi, Luke Zettlemoyer, Percy Liang, Emmanuel Candès, and Tatsunori Hashimoto. S1: Simple test-time scaling. *arXiv preprint arXiv:2405.00002*, 2025.
- [10] OpenAI. Introducing GPT-5. <https://openai.com/blog/introducing-gpt-5>, 2025.
- [11] S. E. Robertson and K. Sparck Jones. Relevance weighting of search terms. *Journal of the American Society for Information Science*, 27(3):129–146, May 1976.
- [12] The rocq-lean-import Development Team. Coq is a lean typechecker. <https://github.com/rocq-community/rocq-lean-import>, 2020.
- [13] Hanzhuo Tan, Qi Luo, Jing Li, and Yuqun Zhang. Llm4decompile: Decompiling binary code with large language models. *arXiv preprint arXiv:2406.00004*, 2024.
- [14] The Coq LSP Development Team. Coq-LSP: Language server protocol implementation for the Coq proof assistant. <https://github.com/ejgallego/coq-lsp>, 2024.
- [15] Laetitia Teodorescu, Guillaume Baudart, Emilio Jesús Gallego Arias, and Marc Lelarge. Nlir: Natural language intermediate representation for mechanized theorem proving. In *MathAI@NeurIPS*, 2024.

155 **A Appendix**

156 **A.1 Training Details**

157 We fine-tune a pretrained and instruction-tuned Qwen2.5-32b-Coder-Instruct[7] for reasoning,
158 following the setup of S1 simple test-time scaling [9]. We train for five epochs with a global batch
159 size of 16, which yields 315 gradient steps. We use bfloat16 precision and an AdamW optimizer
160 with $\beta_1 = 0.9$, $\beta_2 = 0.95$, and weight decay 1×10^{-4} . The learning rate is 1×10^{-5} , warmed up
161 linearly for the first 5% (16 steps) and then cosine-decayed to zero over the remaining 299 steps. We
162 compute loss only on the reasoning traces and the final solutions, and we choose a sequence length
163 large enough to avoid truncation. On GENCI hardware (8×4 NVIDIA H100), this run completes in
164 about one hour. At inference, we sample up to 128 candidates.

165 **A.2 Prompt Babel-translate**

166 For Lean $\rightarrow$ Rocq, we instruct the model to translate a Lean proof term and its local context into a
167 Rocq tactic script. We prepend the following instruction:

```
168 You are given a proof term:
169
170 {term}
171
172 Your task is to derive a sequence of tactics that corresponds to this term.
173
174 When you work through the problem, write down your reasoning in detail inside
175 <think> ... </think> tags. This reasoning should reflect your natural thought
176 process as you explore the structure of the term and figure out what tactics
177 to apply. You should consider different possible approaches, reflect on why
178 some might or might not work, and gradually converge on a tactic choice.
179
180 After each reasoning block, provide the next (group of) tactic(s) enclosed in:
181
182 \box{{
183   <tactic>
184 }}
185
186 Some dependencies that could be helpful:
187
188 {dependencies}
189
190 We hard-code the beginning of the first thinking block:
191
192 <think> Okay, let's try to transform this proof term into a sequence of coq
193 tactics.
194 First let's write down the hypotheses, and the initial goal
195 (after the "|-" symbol) given by the coq proof assistant:
196 {initial_goal}.
```

195 **A.3 Prompt GPT-5 direct translation**

196 For Lean $\rightarrow$ Rocq, we instruct the model to translate a Lean proof script into a Rocq proof script. We
197 give the following instruction:

```
198 You are a precise translator from Lean 4 proofs to Rocq (Coq).
199
200 ## INPUT
201 <ROCQ_STATEMENT>
202 {rocq_statement}
203 </ROCQ_STATEMENT>
204
```

```

205 <LEAN_STATEMENT>
206 {lean_statement}
207 </LEAN_STATEMENT>
208
209 <LEAN_PROOF>
210 {lean_proof}
211 </LEAN_PROOF>
212
213 <DEPENDENCIES>
214 {dependencies}
215 </DEPENDENCIES>
216
217 ## RULES
218 1) Produce two blocks in this exact order and format:
219     a) A thought process.
220     b) The final Rocq proof script delimited by the markers below.
221 2) Use only lemmas/results listed in <DEPENDENCIES>. Do not rely on
222     external lemmas beyond these. Standard Rocq/Coq tactics are allowed.
223 3) Do not invent new lemmas.
224 4) Do not use "Admitted.". End with "Qed.".
225 5) Output must contain only the two blocks below, nothing else.
226
227 ## OUTPUT FORMAT (STRICT)
228 <BEGIN_RATIONALE>
229 [Thought process.]
230 </END_RATIONALE>
231
232 <BEGIN_ROCQ_PROOF>
233 [Place only the Rocq/Coq script here. Inside "Proof." ... "Qed.",
234  do not rewrite the Rocq statement.]
235 </END_ROCQ_PROOF>
236
237 Now produce the output for the given input.

```

In case of an issue, we prepend the following feedback to the next prompt:

```

239 You already tried this proof:
240 {proof}
241 but it fails with this error:
242 {error_message}.

```

243 A.4 Benchmark example

244 We show one Lean/Rocq pair from the aligned set. The Rocq entry is:

```

245 Class AbsField := {
246   R0    : Type;
247   zero  : R0;
248   one   : R0;
249   add   : R0 -> R0 -> R0;
250   mul   : R0 -> R0 -> R0;
251   opp   : R0 -> R0;
252   abs   : R0 -> R0;
253   leR   : R0 -> R0 -> Prop;
254   ltR   : R0 -> R0 -> Prop;
255
256
257   NO     : Type;
258   Nle    : NO -> NO -> Prop;
259   Nmax   : NO -> NO -> NO;

```

```

260 Nle_max_l : forall x y, Nle x (Nmax x y);
261 Nle_max_r : forall x y, Nle y (Nmax x y);
262
263 add_comm : forall x y, add x y = add y x;
264 add_assoc : forall x y z, add (add x y) z = add x (add y z);
265 add_zero : forall x, add x zero = x;
266 add_opp : forall x, add x (opp x) = zero;
267 mul_comm : forall x y, mul x y = mul y x;
268 mul_assoc : forall x y z, mul (mul x y) z = mul x (mul y z);
269 mul_one : forall x, mul x one = x;
270
271 le_refl : forall x, leR x x;
272 le_trans : forall x y z, leR x y -> leR y z -> leR x z;
273 add_le_add : forall a b c d, leR a b -> leR c d -> leR (add a c)\\
274 (add b d);
275
276 abs_nonneg : forall x, leR zero (abs x);
277 abs_triangle : forall x y, leR (abs (add x y)) (add (abs x) (abs y));
278 abs_opp : forall x, abs (opp x) = abs x;
279 abs_sub_symm : forall x y, abs (add x (opp y)) = abs (add y (opp x));
280 sub_decomp : forall x y z, add x (opp z) = add (add x (opp y)) \\
281 (add y (opp z));
282 sub_eq_zero : forall x y, add x (opp y) = zero -> x = y;
283
284
285 eq_of_forall_eps2 : forall x, (forall eps, ltR zero eps -> leR \\
286 (abs x) (add eps eps)) -> x = zero
287 }.
288
289 Section Limits.
290 Context {A : AbsField}.
291
292 Definition sub (x y : R0) : R0 := add x (opp y).
293
294 Definition limit (u : N0 -> R0) (l : R0) : Prop :=
295   forall eps, ltR zero eps -> exists N, forall n, Nle N n -> leR\\
296   (abs (sub (u n) l)) eps.
297
298 Lemma abs_sub_triangle (x y z : R0) :
299   leR (abs (sub x z)) (add (abs (sub x y)) (abs (sub y z))).
300 Proof.
301   unfold sub.
302   rewrite (sub_decomp x y z).
303   apply abs_triangle.
304 Qed.
305
306 Theorem limit_unique (u : N0 -> R0) (l m : R0) :
307   limit u l -> limit u m -> l = m.
308 Proof.
309   intros Hl Hm.
310   assert (Hbound: forall eps, ltR zero eps -> leR (abs (sub l m))\\
311   (add eps eps)).
312   { intros eps Heps.
313     destruct (Hl eps Heps) as [N1 HN1].
314     destruct (Hm eps Heps) as [N2 HN2].
315     set (N := Nmax N1 N2).
316     assert (H1u: leR (abs (sub (u N) l)) eps).
317     { apply HN1. apply Nle_max_l. }
318     assert (H1: leR (abs (sub l (u N))) eps).

```



```

319     { unfold sub in H1u. rewrite abs_sub_symm in H1u. fold sub\\
320       in H1u. exact H1u. }
321     assert (H2: leR (abs (sub (u N) m)) eps).
322     { apply HN2. apply N1e_max_r. }
323     assert (Htri: leR (abs (sub 1 m)) (add (abs (sub 1 (u N))))\\
324       (abs (sub (u N) m)))) by apply abs_sub_triangle.
325     eapply le_trans; [exact Htri|].
326     apply add_le_add; [exact H1|exact H2].
327   }
328   assert (Hz: sub 1 m = zero).
329   { apply eq_of_forall_eps2. exact Hbound. }
330   unfold sub in Hz.
331   now apply sub_eq_zero in Hz.
332   Qed.
333
334   End Limits.

335   The Lean counterpart is:

336   class AbsField (R : Type) where
337     zero  : R
338     one   : R
339     add   : R → R → R
340     mul   : R → R → R
341     opp   : R → R
342     abs   : R → R
343     leR   : R → R → Prop
344     ltR   : R → R → Prop
345
346
347     NatAlt      : Type
348     NatAltle    : NatAlt → NatAlt → Prop
349     NatAltMax   : NatAlt → NatAlt → NatAlt
350     le_max_left : forall x y, NatAltle x (NatAltMax x y)
351     le_max_right : forall x y, NatAltle y (NatAltMax x y)
352
353     add_comm    : forall x y, add x y = add y x
354     add_assoc   : forall x y z, add (add x y) z = add x (add y z)
355     add_zero    : forall x, add x zero = x
356     add_opp     : forall x, add x (opp x) = zero
357     opp_add     : forall x y, opp (add x y) = add (opp x) (opp y)
358     mul_comm    : forall x y, mul x y = mul y x
359     mul_assoc   : forall x y z, mul (mul x y) z = mul x (mul y z)
360     mul_one     : forall x, mul x one = x
361
362     le_refl     : forall x, leR x x
363     le_trans    : forall x y z, leR x y → leR y z → leR x z
364     add_le_add  : forall a b c d, leR a b → leR c d → leR (add a c) \\
365       (add b d)
366
367     abs_nonneg  : forall x, leR zero (abs x)
368     abs_triangle : forall x y, leR (abs (add x y)) (add (abs x) (abs y))
369     abs_opp     : forall x, abs (opp x) = abs x
370     abs_sub_symm : forall x y, abs (add x (opp y)) = abs (add y (opp x))
371
372
373     sub_decomp  : forall x y z, add x (opp z) = add (add x (opp y)) \\
374       (add y (opp z))
375     sub_eq_zero : forall x y, add x (opp y) = zero → x = y
376

```

```

377   eq_of_forall_eps2 : forall x, (forall eps, ltR zero eps → leR (abs x) \\  

378     (add eps eps)) → x = zero  

379  

380 namespace Limits  

381  

382 variable {R : Type} [AR : AbsField R]  

383  

384 def sub (x y : R) : R := AbsField.add x (AbsField.opp y)  

385  

386 def limit (u : (AbsField.NatAlt (R := R)) → R) (l : R) : Prop :=  

387   forall eps : R, AbsField.ltR AbsField.zero eps →  

388     exists N : AbsField.NatAlt (R := R),  

389     forall n : AbsField.NatAlt (R := R),  

390       AbsField.NatAltle (R := R) N n →  

391       AbsField.leR (AbsField.abs (sub (u n) l)) eps  

392  

393 theorem abs_sub_triangle (x y z : R) :  

394   AbsField.leR (AbsField.abs (sub x z)) (AbsField.add (AbsField.abs \\  

395     (sub x y)) (AbsField.abs (sub y z))) := by  

396  

397   have : sub x z = AbsField.add (sub x y) (sub y z) := sub_decomp x y z  

398   simpa [this] using AbsField.abs_triangle (sub x y) (sub y z)  

399  

400 theorem limit_unique (u : (AbsField.NatAlt (R := R)) → R) (l m : R) :  

401   limit u l → limit u m → l = m :=  

402   by  

403     intro Hl Hm  

404  

405     have Hbound' : forall eps, AR.ltR AR.zero eps → AR.leR (AR.abs (sub l m)) \\  

406       (AR.add eps eps) := by  

407       intro eps Heps  

408       rcases Hl eps Heps with ⟨N1, HN1⟩  

409       rcases Hm eps Heps with ⟨N2, HN2⟩  

410       let N := AbsField.NatAltMax (R := R) N1 N2  

411       have H1 : AbsField.leR (AbsField.abs (sub (u N) l)) eps := by  

412         apply HN1  

413         exact AbsField.le_max_left (R := R) _ _  

414       have H2 : AbsField.leR (AbsField.abs (sub (u N) m)) eps := by  

415         apply HN2  

416         exact AbsField.le_max_right (R := R) _ _  

417  

418       have Htri : AbsField.leR (AbsField.abs (sub l m)) (AbsField.add \\  

419         (AbsField.abs (sub l (u N))) (AbsField.abs (sub (u N) m))) := by  

420         simpa using abs_sub_triangle l (u N) m  

421  

422  

423       have H1' : AbsField.leR (AbsField.abs (sub l (u N))) eps := by  

424  

425         simpa [sub, AbsField.abs_sub_symm (u N) l] using H1  

426         exact AbsField.le_trans _ _ _ Htri (AbsField.add_le_add _ _ _ H1' H2)  

427  

428       have Hz : sub l m = AbsField.zero := AbsField.eq_of_forall_eps2 (sub l m) Hbound'  

429  

430       exact AbsField.sub_eq_zero l m (by simpa [sub] using Hz)  

431 end Limits

```

432 A.5 Translation to SSReflect

433 We illustrate the Rocq $\rightarrow$ SSReflect translation on a lemma extracted from CoRN. Babel takes as
434 input the proof term and produces a short SSReflect script that closes the goal under the same context.

435 Lemma (CoRN).

```
436 Lemma CR_nonZero_as_Cauchy_IR_nonZero_1 :
437   forall (x:CR), (CRapartT x 0)%CR ->
438   Dom (f_rcpcl' _) (CRasCauchy_IR x).
```

439 Proof term.

```
440 CR_nonZero_as_Cauchy_IR_nonZero_1 =
441 fun (x : CR) (x_ : (x >< ' 0%Q)%CR) =>
442   ap_wdr_unfolded Cauchy_IR (CRasCauchy_IR x) (CRasCauchy_IR (' 0%Q)%CR) [0]
443   (CR_ap_as_Cauchy_IR_ap_1 x (' 0%Q)%CR x_)
444   (CR_inject_Q_as_Cauchy_IR_inject_Q 0)
445 : forall x : CR,
446   (x >< ' 0%Q)%CR ->
447   Dom (f_rcpcl' Cauchy_IR) (CRasCauchy_IR x)
448
449 Arguments CR_nonZero_as_Cauchy_IR_nonZero_1 x _
```

450 Original Rocq script.

```
451 intros x x_.
452 change ((CRasCauchy_IR x)[#][0]).
453 stepr (CRasCauchy_IR 0%CR).
454 apply CR_ap_as_Cauchy_IR_ap_1.
455 assumption.
456 eapply CR_inject_Q_as_Cauchy_IR_inject_Q.
457 Defined.
```

458 Babel (SSReflect) script from the term.

```
459 move=> x_; rewrite /CRapartT; apply: CR_ap_as_Cauchy_IR_ap_1 => /=.
```

460 A.6 Translation from Lean to Rocq

461 We illustrate the Lean $\rightarrow$ Rocq translation on a lemma whose Lean proof uses the `simp` tactic. Babel
462 takes as input the proof term and the goal (it does not see the Lean proof), and reconstructs the
463 explicit calls to lemmas such as `map_smul` and `smul_zeroV` that `simp` applies implicitly.

464 Lemma (Lean).

```
465 theorem ker_smul {L : LMap (R:=R) V W} {a : R} {x : V} :
466   ker L x  $\rightarrow$  ker L (VSpace.smul (R:=R) a x) :=
467   by
468     intro hx
469     unfold ker at *
470     simp [hx]
```

471 Proof term.

```
472 ker_smul =
473 fun (VS2 : VSpace) (L : Lin) (a : R0) (x : V0) (hx : ker L x) =>
474   eq_ind_r (fun v : V0 => v = zeroV)
475     (eq_ind_r (fun v : V0 => (a •V v) = zeroV)
476       (eq_ind_r (fun v : V0 => v = zeroV) eq_refl (smul_zeroV a)) hx)
```

```

477     (map_smul L a x)
478   : forall {VS2 : VSpace} (L : Lin) (a : RO) (x : VO),
479     ker L x -> ker L (a •V x)

```

480 **Original Lean script.**

```

481 intro hx
482 unfold ker at *
483 simp [hx]

```

484 **Babel (Rocq) reconstruction.** The beginning of the first reasoning block (initial hypotheses and goal) is hardcoded by us to provide context; it is *not* generated by the model.

486 Okay, let's try to transform this proof term into a sequence of Coq tactics.
487 First, let's write down the hypotheses and the initial goal (after the "|-")
488 given by the Coq proof assistant:

```

489
490 F : Field
491 VS1, VS2 : VSpace
492 L : Lin
493 a : RO
494 x : VO
495 |- ker L x -> ker L (a •V x)

```

496

```

497 The proof term starts with fun {R} ... {x} hx => .... This is a lambda
498 abstraction that binds the variable hx. In a Coq proof, this corresponds
499 to introducing an assumption. The type of this assumption is the premise
500 of the implication in the goal, which is ker L x. So, the first logical
501 step is to introduce this assumption into our proof context. The tactic
502 for this is intros. The lambda term gives it the name hx, so I'll use
503 that. This will change our goal to the conclusion of the implication,
504 which is ker L (a •V x). </think>
505 \box{intros hx.}
506
507 <think> Now, my proof state is:
508 F, VS1, VS2, L, a, x : ...
509 hx : ker L x
510 |- ker L (a •V x)
511
512 The goal is ker L (a •V x). This is a definition. To make progress, I
513 need to understand what I'm trying to prove. ker is the kernel of a
514 linear map, which means it's the set of vectors that the map sends to zero.
515 So, ker L (a •V x) should mean L (a •V x) = 0. To make the proof term
516 easier to work with, I should unfold this definition. The tactic for
517 unfolding definitions in the goal is unfold. Let's apply that to ker in the
518 goal. </think>
519 \box{unfold ker in |- *.}
520
521 <think> The goal has now unfolded to L (a •V x) = Lin.VSpace.zeroV R. The
522 lambda term contains L.map_smul a x. This looks like a lemma about linear
523 maps and scalar multiplication. It probably states that L (a •V x) =
524 a •V (L x). This seems like a very useful fact, as it lets me distribute the
525 linear map L over the scalar multiplication •V. My current goal's left-hand
526 side is L (a •V x), which is the perfect form to apply this lemma. The tactic
527 rewrite is used to apply an equality to the goal. Since map_smul is a lemma
528 of the form LHS = RHS, rewrite map_smul will replace LHS with RHS. I expect
529 this to change my goal to a •V (L x) = Lin.VSpace.zeroV R. This looks like a
530 promising simplification. </think>

```

```

531 \box{rewrite map_smul.}
532
533 <think> My goal is now a  $\bullet V (L x) = \text{Lin.VSpace.zeroV } R$ . Let's look at my
534 hypotheses. I have  $hx : \ker L x$ . This is still using the  $\ker$  definition.
535 It would be helpful to have this hypothesis in its unfolded form as well,
536  $L x = \text{Lin.VSpace.zeroV } R$ . This will make it easier to use  $hx$  to rewrite parts
537 of my goal. I'll use the unfold tactic again, this time on the hypothesis  $hx$ .
538 </think>
539 \box{unfold ker in hx.}
540
541 <think> Now my context contains  $hx : L x = \text{Lin.VSpace.zeroV } R$ , and my goal is
542 a  $\bullet V (L x) = \text{Lin.VSpace.zeroV } R$ . This is great. The term  $L x$  in my goal appears
543 exactly as the left-hand side of the equality in my hypothesis  $hx$ . This is a
544 perfect opportunity to use the rewrite tactic. By doing rewrite hx, I can
545 replace  $L x$  in the goal with  $\text{Lin.VSpace.zeroV } R$ . I expect this to transform
546 my goal into a  $\bullet V (\text{Lin.VSpace.zeroV } R) = \text{Lin.VSpace.zeroV } R$ . </think>
547 \box{rewrite hx.}
548
549 <think> The goal is now a  $\bullet V (\text{Lin.VSpace.zeroV } R) = \text{Lin.VSpace.zeroV } R$ . Looking
550 at the lambda term, I see a reference to Lin.VSpace.smul_zero a. This must be
551 the lemma for  $a \bullet V 0 = 0$ . My goal's left-hand side, a  $\bullet V (\text{Lin.VSpace.zeroV } R)$ ,
552 is precisely the form required by this lemma. Once again, rewrite is the tactic
553 to use. Applying rewrite smul_zeroV should simplify the left-hand side of my
554 goal to  $\text{Lin.VSpace.zeroV } R$ , making the
555 goal  $\text{Lin.VSpace.zeroV } R = \text{Lin.VSpace.zeroV } R$ . </think>
556 \box{rewrite smul_zeroV.}
557
558 <think> After the last rewrite, my goal is  $\text{Lin.VSpace.zeroV } R = \text{Lin.VSpace.zeroV } R$ .
559 This is an equality of a term with itself. This is a trivially true statement. There
560 are several tactics to solve such a goal. reflexivity is the most direct one, as it
561 is designed specifically for proving  $t = t$ . The lambda term structure Eq.trans ...
562 (eq_self ...) confirms that the proof concludes by showing both sides are the same
563 thing. Therefore, reflexivity is the final step to close the proof. </think>
564 \box{reflexivity.}

```